

PQHS471 HW 2

Gregory Powers (gxp145)

Chapter 5 Q9

(A)

```
library(MASS)
library(Matrix)
library(knitr)
library(kableExtra)
library(MVN)
library(corrplot)
attach(Boston)
set.seed(1)
medv.Mean <- mean(medv)
medv.Mean
```

```
## [1] 22.53281
```

(B) As I am learning R, I'm going to do this a few ways.

```
sum(medv > 0)
```

```
## [1] 506
```

```
length(medv)
```

```
## [1] 506
```

```
medv.Error <- sd(medv)/sqrt(506)
medv.Error
```

```
## [1] 0.4088611
```

```
print(sd(medv)/sqrt(length(medv)))
```

```
## [1] 0.4088611
```

(C)

```
library(boot)
mean.fn <- function (x ,id) {
  return(mean(x[id]))
}
boot.M <- boot(medv, mean.fn, 1000)
boot.M
```

```
##
## ORDINARY NONPARAMETRIC BOOTSTRAP
##
##
## Call:
## boot(data = medv, statistic = mean.fn, R = 1000)
##
##
## Bootstrap Statistics :
```

```
##      original      bias      std. error
## t1* 22.53281 0.008517589  0.4119374
```

```
boot.SD <- sd(boot.M$t)/length(t)
boot.SD - medv.Error
```

```
## [1] 0.003076292
```

The difference between the bootstrapped estimate and the original is about 0.01.

(D)

```
c(boot.M$t0 - 2 *sd(boot.M$t), boot.M$t0 + 2 *sd(boot.M$t))
```

```
## [1] 21.70893 23.35668
```

```
t.test(medv)
```

```
##
## One Sample t-test
##
## data: medv
## t = 55.111, df = 505, p-value < 2.2e-16
## alternative hypothesis: true mean is not equal to 0
## 95 percent confidence interval:
## 21.72953 23.33608
## sample estimates:
## mean of x
## 22.53281
```

(E)

```
medv.Median <- median(Boston$medv); medv.Median
```

```
## [1] 21.2
```

(F)

```
median.fn <- function (x ,id) {
  return(median(x[id]))
}

boot.Median <- boot(medv, median.fn, 1000)
boot.Median
```

```
##
## ORDINARY NONPARAMETRIC BOOTSTRAP
##
##
## Call:
## boot(data = medv, statistic = median.fn, R = 1000)
##
##
## Bootstrap Statistics :
##      original      bias      std. error
## t1*      21.2 -0.0098  0.3874004
```

The estimated SE of the median is 0.3801

(G)

```
print(medv.muTen <- quantile(medv, 0.1))

## 10%
## 12.75

quantile.fn <- function(x ,id) {
  return(quantile(x[id], 0.1))
}

boot.Quantile10 <- boot(medv, quantile.fn, 1000)
boot.Quantile10

##
## ORDINARY NONPARAMETRIC BOOTSTRAP
##
##
## Call:
## boot(data = medv, statistic = quantile.fn, R = 1000)
##
##
## Bootstrap Statistics :
##      original  bias    std. error
## t1*      12.75 0.00515    0.5113487

The estimated SE is 0.4826.
```

Chapter 6 Q9

(A)

```
library(ISLR)
detach(Boston)
attach(College)
```

```
library(dplyr)
summary(College)
```

```
## Private      Apps      Accept      Enroll      Top10perc
## No :212      Min.   : 81      Min.   : 72      Min.   : 35      Min.   : 1.00
## Yes:565      1st Qu.: 776      1st Qu.: 604      1st Qu.: 242      1st Qu.:15.00
##              Median : 1558      Median : 1110      Median : 434      Median :23.00
##              Mean   : 3002      Mean   : 2019      Mean   : 780      Mean   :27.56
##              3rd Qu.: 3624      3rd Qu.: 2424      3rd Qu.: 902      3rd Qu.:35.00
##              Max.   :48094      Max.   :26330      Max.   :6392      Max.   :96.00
## Top25perc    F.Undergrad  P.Undergrad    Outstate
## Min.   : 9.0      Min.   : 139      Min.   : 1.0      Min.   : 2340
## 1st Qu.: 41.0      1st Qu.: 992      1st Qu.: 95.0      1st Qu.: 7320
## Median : 54.0      Median : 1707      Median : 353.0      Median : 9990
## Mean   : 55.8      Mean   : 3700      Mean   : 855.3      Mean   :10441
## 3rd Qu.: 69.0      3rd Qu.: 4005      3rd Qu.: 967.0      3rd Qu.:12925
## Max.   :100.0      Max.   :31643      Max.   :21836.0      Max.   :21700
## Room.Board   Books      Personal      PhD
## Min.   :1780      Min.   : 96.0      Min.   : 250      Min.   : 8.00
## 1st Qu.:3597      1st Qu.: 470.0      1st Qu.: 850      1st Qu.: 62.00
```

```
## Median :4200    Median : 500.0    Median :1200    Median : 75.00
## Mean   :4358    Mean   : 549.4    Mean   :1341    Mean   : 72.66
## 3rd Qu.:5050    3rd Qu.: 600.0    3rd Qu.:1700    3rd Qu.: 85.00
## Max.   :8124    Max.   :2340.0    Max.   :6800    Max.   :103.00
##      Terminal      S.F.Ratio      perc.alumni      Expend
## Min.    : 24.0    Min.    : 2.50    Min.    : 0.00    Min.    : 3186
## 1st Qu.: 71.0    1st Qu.:11.50    1st Qu.:13.00    1st Qu.: 6751
## Median : 82.0    Median :13.60    Median :21.00    Median : 8377
## Mean   : 79.7    Mean   :14.09    Mean   :22.74    Mean   : 9660
## 3rd Qu.: 92.0    3rd Qu.:16.50    3rd Qu.:31.00    3rd Qu.:10830
## Max.   :100.0    Max.   :39.80    Max.   :64.00    Max.   :56233
##      Grad.Rate
## Min.    : 10.00
## 1st Qu.: 53.00
## Median : 65.00
## Mean   : 65.46
## 3rd Qu.: 78.00
## Max.   :118.00
```

```
str(College)
```

```
## 'data.frame':    777 obs. of  18 variables:
## $ Private      : Factor w/ 2 levels "No","Yes": 2 2 2 2 2 2 2 2 2 2 ...
## $ Apps         : num  1660 2186 1428 417 193 ...
## $ Accept       : num  1232 1924 1097 349 146 ...
## $ Enroll       : num  721 512 336 137 55 158 103 489 227 172 ...
## $ Top10perc    : num  23 16 22 60 16 38 17 37 30 21 ...
## $ Top25perc    : num  52 29 50 89 44 62 45 68 63 44 ...
## $ F.Undergrad  : num  2885 2683 1036 510 249 ...
## $ P.Undergrad  : num  537 1227 99 63 869 ...
## $ Outstate     : num  7440 12280 11250 12960 7560 ...
## $ Room.Board   : num  3300 6450 3750 5450 4120 ...
## $ Books        : num  450 750 400 450 800 500 500 450 300 660 ...
## $ Personal     : num  2200 1500 1165 875 1500 ...
## $ PhD          : num  70 29 53 92 76 67 90 89 79 40 ...
## $ Terminal     : num  78 30 66 97 72 73 93 100 84 41 ...
## $ S.F.Ratio    : num  18.1 12.2 12.9 7.7 11.9 9.4 11.5 13.7 11.3 11.5 ...
## $ perc.alumni  : num  12 16 30 37 2 11 26 37 23 15 ...
## $ Expend       : num  7041 10527 8735 19016 10922 ...
## $ Grad.Rate    : num  60 56 54 59 15 55 63 73 80 52 ...
```

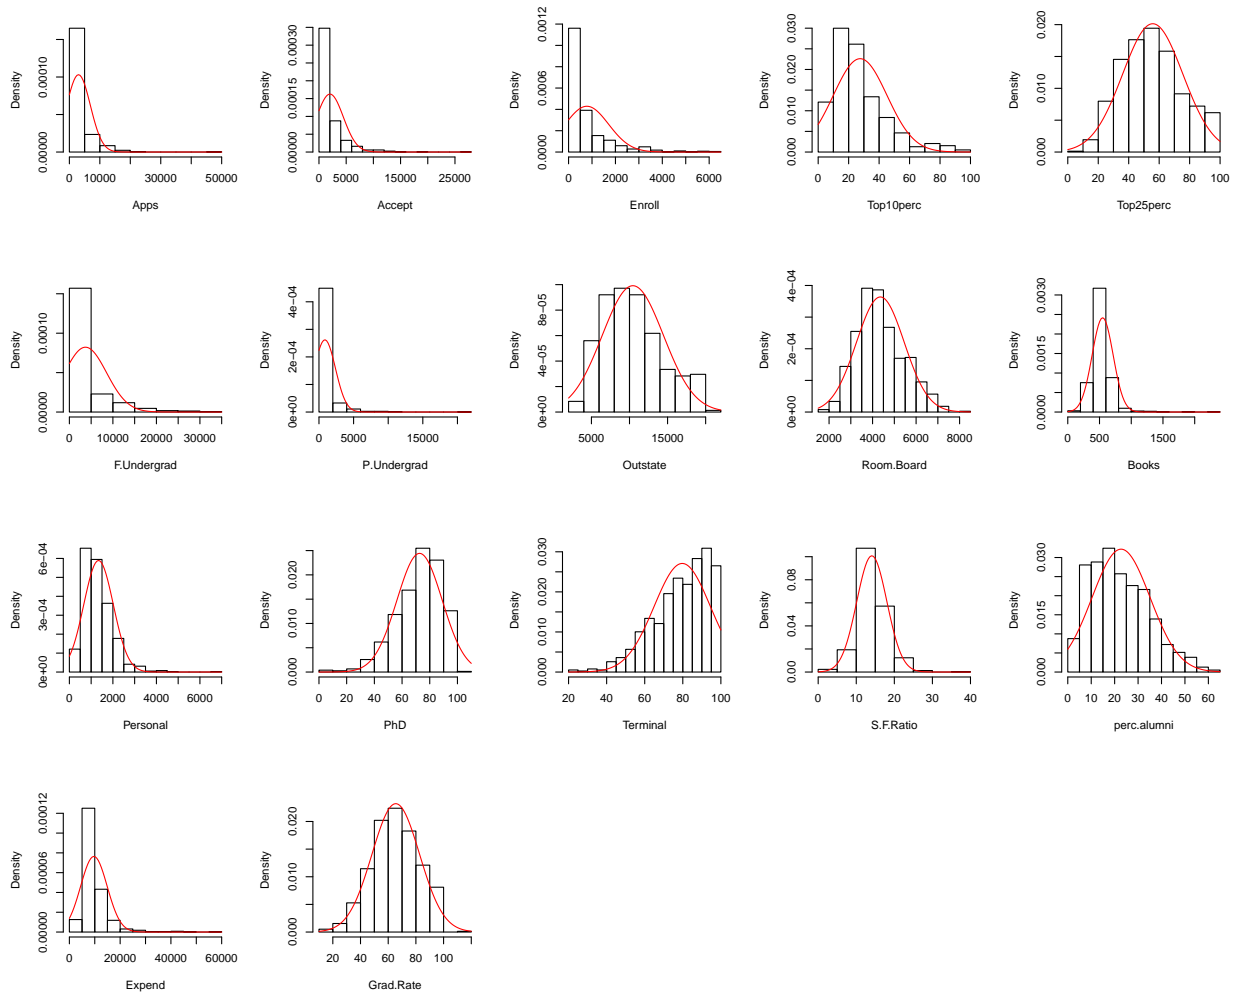
```
anyDuplicated(College)
```

```
## [1] 0
```


```
sum(is.na(College))
```

```
## [1] 0
```

```
uniPlot(College[2:18], type = "histogram")
```



```
r <- cor(College[2:18])
title <- 'College Data'
corrplot(r, method = "color", type = 'upper', diag = FALSE, addCoef.col = "black",
         order = "hclust", title = title, mar=c(0,0,1,0))
```



Using dplyr to split into test and train.

```
col.train <- sample_frac(College, 0.8)
col.test = setdiff(College, col.train)
nrow(col.train) + nrow(col.test) == nrow(College)
```

```
## [1] TRUE
```

(B)

```
lm.Apps <- lm(Apps ~ ., data = col.train)
summary(lm.Apps)
```

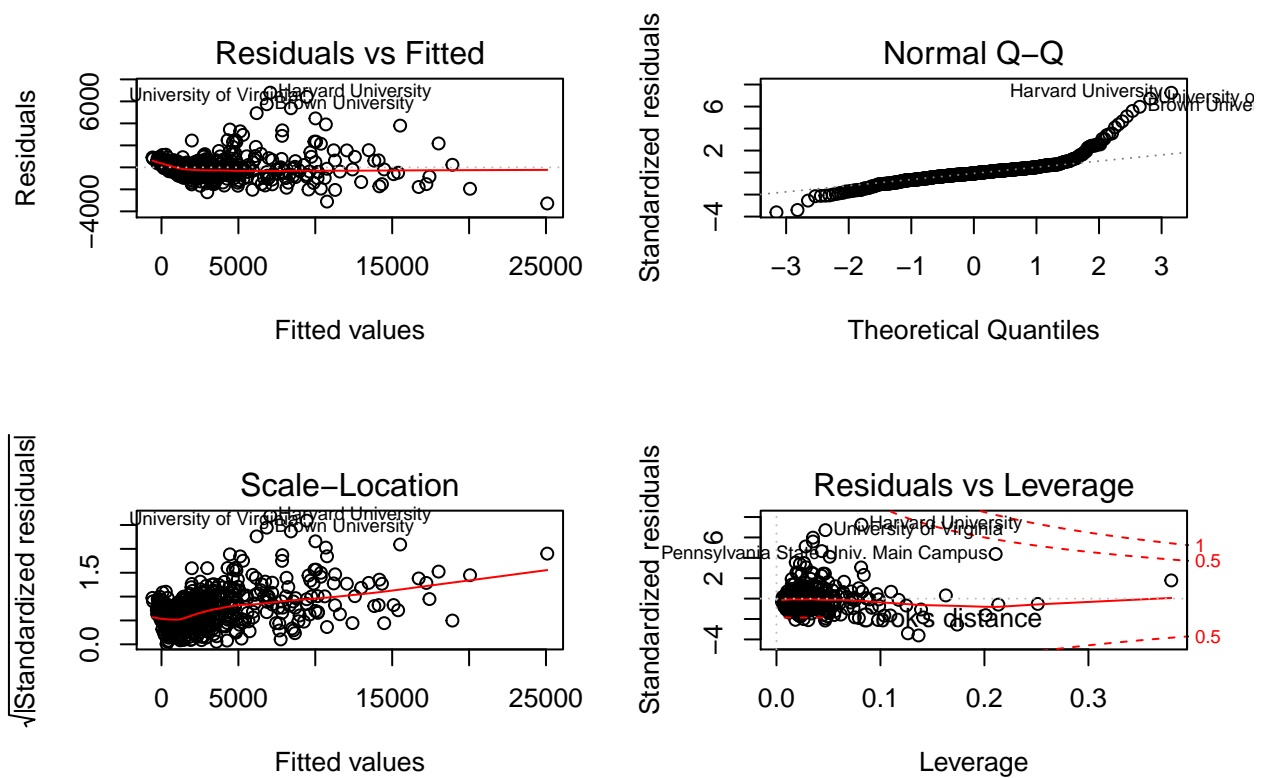
```
##
## Call:
## lm(formula = Apps ~ ., data = col.train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3280.5  -436.5   -72.1    291.8   6792.1
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
```

```

## (Intercept) -4.573e+02  4.208e+02  -1.087  0.277607
## PrivateYes  -6.156e+02  1.470e+02  -4.187  3.24e-05 ***
## Accept       1.309e+00  5.261e-02  24.878  < 2e-16 ***
## Enroll      -4.010e-01  2.045e-01  -1.961  0.050346 .
## Top10perc    3.967e+01  6.017e+00   6.592  9.46e-11 ***
## Top25perc   -9.843e+00  4.846e+00  -2.031  0.042688 *
## F.Undergrad  8.041e-02  3.461e-02   2.323  0.020501 *
## P.Undergrad  3.213e-02  3.148e-02   1.021  0.307793
## Outstate    -6.510e-02  2.054e-02  -3.170  0.001603 **
## Room.Board   1.777e-01  5.032e-02   3.531  0.000446 ***
## Books        5.272e-03  2.436e-01   0.022  0.982742
## Personal    -1.720e-02  6.511e-02  -0.264  0.791696
## PhD         -6.635e+00  5.033e+00  -1.318  0.187896
## Terminal    -5.347e+00  5.466e+00  -0.978  0.328344
## S.F.Ratio    1.980e+00  1.409e+01   0.141  0.888253
## perc.alumni -7.617e+00  4.308e+00  -1.768  0.077519 .
## Expend       8.415e-02  1.261e-02   6.672  5.72e-11 ***
## Grad.Rate    1.310e+01  3.268e+00   4.009  6.87e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 977.5 on 604 degrees of freedom
## Multiple R-squared:  0.9292, Adjusted R-squared:  0.9272
## F-statistic:  466 on 17 and 604 DF,  p-value: < 2.2e-16

par(mfrow=c(2,2))
plot(lm.Apps)

```




```
lm.pred <- predict(lm.Apps, col.test)
lm.er <- mean((col.test[, 'Apps'] - lm.pred)^2) #doing it the book way once. Referencing cols by [] is
```

Test RSS = 944829

(C)

```
library(glmnet)
library(foreach)
train.mat <- model.matrix(Apps ~ ., data = col.train) # Reminder to self: no missing [, -1]
test.mat <- model.matrix(Apps ~ ., data = col.test) # [, -1]
grid = 10^seq(10, -2, length = 100)
ridge.cv <- cv.glmnet(train.mat, col.train$Apps, alpha = 0, lambda = grid, thresh = 1e-12) # can also r
plot(ridge.cv)
```

```
ridge.cv$lambda.min
```

```
## [1] 14.17474
```


```
ridge.pred <- predict(ridge.cv, newx = test.mat, s = ridge.cv$lambda.min)
```

```
ridge.er <- mean((col.test$Apps - ridge.pred)^2)
```

The RSS improves only slightly.

(D)

```
lass.cv <- cv.glmnet(train.mat, col.train$Apps, alpha = 1, lambda = grid, thresh = 1e-12)
plot(lass.cv)
```




```
lass.pred <- predict(lass.cv, s = lass.cv$lambda.min, newx = test.mat)
lass.er <- mean((col.test$Apps - lass.pred)^2)
mod.lass = glmnet(model.matrix(Apps~., data=College), College$Apps, alpha=1)
mod.lass.p <-predict(mod.lass, s=lass.cv$lambda.min, type="coefficients"); mod.lass.p
```

```
## 19 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept) -556.01427783
## (Intercept) .
## PrivateYes  -458.50058120
## Accept      1.49947647
## Enroll      -0.34285887
## Top10perc   39.10796866
## Top25perc   -6.17240648
## F.Undergrad .
## P.Undergrad 0.03388810
## Outstate    -0.06812801
## Room.Board  0.13486936
## Books       .
## Personal    0.01173185
## PhD         -6.71936715
## Terminal    -3.15826314
## S.F.Ratio    8.61896397
## perc.alumni -0.72038115
## Expend       0.07249557
## Grad.Rate    6.28499717
```

The test RSS is 2136982. There are 14 nonzero coefficient estimates, though some are quite small.

(E)

```
library(pls)
pcr.mod <- pcr(Apps ~ ., data = col.train, scale = TRUE, validation = "CV")
validationplot(pcr.mod, val.type = 'MSEP')
```



```
summary(pcr.mod)
```

```
## Data:      X dimension: 622 17
## Y dimension: 622 1
## Fit method: svdpc
## Number of components considered: 17
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              3625    3485    1720    1729    1373    1323    1294
## adjCV           3625    3484    1718    1730    1327    1307    1291
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV          1264    1213    1199    1198    1199    1200    1204
## adjCV        1263    1209    1196    1196    1197    1197    1201
##      14 comps 15 comps 16 comps 17 comps
## CV           1204    1206    1028    1019
## adjCV        1201    1203    1025    1016
##
## TRAINING: % variance explained
```

```
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps
## X      31.540   57.13   64.18   69.91   75.36   80.53   84.52
## Apps    7.907   77.82   77.90   87.08   87.61   87.77   88.38
##      8 comps  9 comps 10 comps 11 comps 12 comps 13 comps 14 comps
## X      88.04   90.96   93.17   95.27   97.07   98.16   98.94
## Apps    89.24   89.70   89.74   89.82   89.86   89.86   89.88
##     15 comps 16 comps 17 comps
## X      99.45   99.86  100.00
## Apps    90.17   92.60   92.92
```

```
pcr.pred <- predict(pcr.mod, col.test, ncomp = 17)
pcr.er <- mean((col.test$Apps - pcr.pred)^2)
```


M = 17 (no reduction in dimensions), test RSS = 1969505

```
pls.mod <- plsrf(Apps ~ ., data = col.train, scale = TRUE, validation = "CV")
summary(pls.mod)
```

```
## Data:      X dimension: 622 17
## Y dimension: 622 1
## Fit method: kernelppls
## Number of components considered: 17
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV           3625   1553   1329   1161   1122   1065   1043
## adjCV        3625   1551   1332   1159   1118   1056   1038
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV           1029   1025   1023   1020   1020   1021   1021
## adjCV        1026   1022   1021   1017   1017   1018   1018
##     14 comps 15 comps 16 comps 17 comps
## CV           1020   1020   1020   1020
## adjCV        1017   1017   1017   1017
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps
## X      26.32   47.39   62.60   65.86   68.51   73.43   76.87
## Apps    81.95   86.86   90.23   91.21   92.35   92.71   92.80
##      8 comps  9 comps 10 comps 11 comps 12 comps 13 comps 14 comps
## X      80.61   83.59   85.50   88.59   91.60   94.26   96.03
## Apps    92.83   92.86   92.89   92.90   92.91   92.91   92.92
##     15 comps 16 comps 17 comps
## X      96.87   98.81  100.00
## Apps    92.92   92.92   92.92
```

```
validationplot(pls.mod, val.type = 'MSEP')
```


Apps



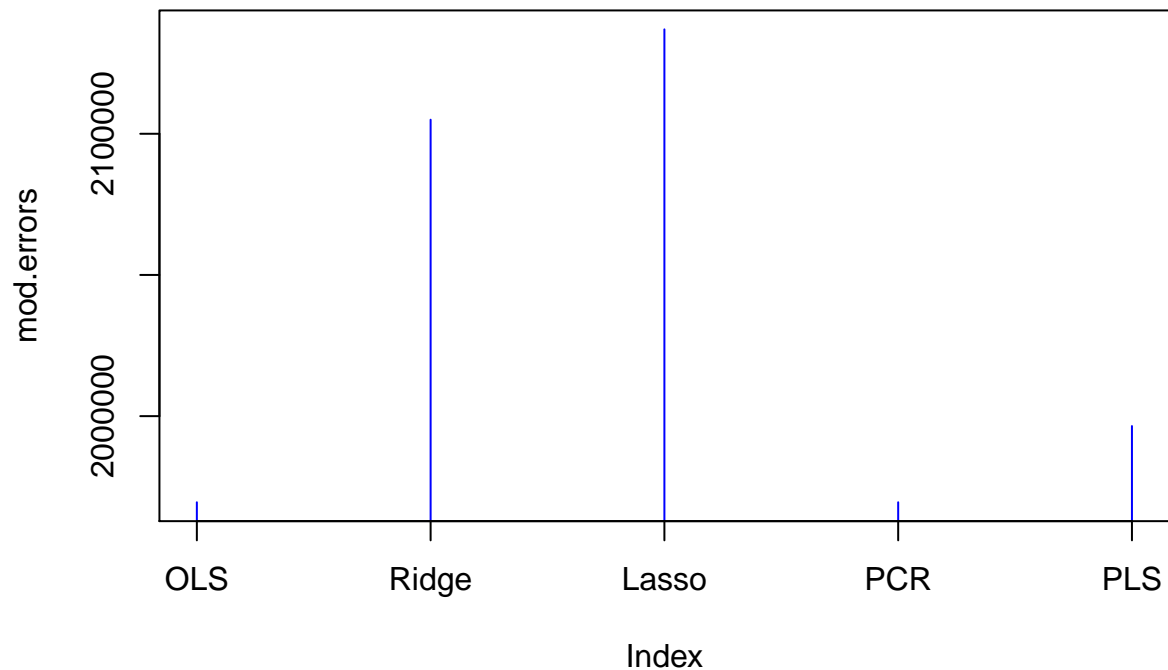
```
pls.pred <- predict(pls.mod, col.test, ncomp = 10)
pls.er <- mean((col.test$Apps - pls.pred)^2)
```

Min CV is where M = 10. Test error RSS = 1996495

```
mod.errors <- c(lm.er, ridge.er, lass.er, pcr.er, pls.er)
names(mod.errors) <- c('OLS', 'Ridge', 'Lasso', 'PCR', 'PLS')
barplot(mod.errors, main = 'Test Error by Method', xlab = 'Method', ylab = 'RSS')
```



```
plot(mod.errors, type = 'h', col='blue', xaxt='n')  
axis(1, at=1:5, lab=c('OLS', 'Ridge', 'Lasso', 'PCR', 'PLS'))
```




```
kable(mod.errors, col.names = 'RSS')%>%
  kable_styling()
```

| | RSS |
|-------|---------|
| OLS | 1969505 |
| Ridge | 2104991 |
| Lasso | 2136982 |
| PCR | 1969505 |
| PLS | 1996495 |

```
#The idea to do the below came from an R forum
t.avg <- mean(col.test$Apps)
ols.r2 = 1 - mean((lm.pred - col.test$Apps)^2) / mean((t.avg - col.test$Apps)^2)
ridge.r2 = 1 - mean((ridge.pred - col.test$Apps)^2) / mean((t.avg - col.test$Apps)^2)
lass.r2 = 1 - mean((lass.pred - col.test$Apps)^2) / mean((t.avg - col.test$Apps)^2)
pcr.r2 = 1 - mean((pcr.pred - col.test$Apps)^2) / mean((t.avg - col.test$Apps)^2)
pls.r2 = 1 - mean((pls.pred - col.test$Apps)^2) / mean((t.avg - col.test$Apps)^2)

mod.effect <- c(ols.r2, ridge.r2, lass.r2, pcr.r2, pls.r2)
names(mod.effect) <- c('OLS', 'Ridge', 'Lasso', 'PCR', 'PLS')

plot(mod.effect, type = 'h', col='blue', xaxt='n')
axis(1, at=1:5, lab=c('OLS', 'Ridge', 'Lasso', 'PCR', 'PLS'))
```



```
kable(mod.effect, col.names = 'r2')%>%
  kable_styling()
```

| | r2 |
|-------|-----------|
| OLS | 0.9122158 |
| Ridge | 0.9061770 |
| Lasso | 0.9047511 |
| PCR | 0.9122158 |
| PLS | 0.9110128 |

There is very little difference in RSS and all models account for the variation in applications quite well ($r^2 < .9$). OLS followed by PLS produce the models with the smallest test error and largest r^2 , though the absolute differences are tiny. That PCR is not considered: as no dimensions were reduced, it is equivalent to OLS.

Chapter 6 Q11

(A)

```
#Splitting into test and train via dplyr. Making test Mats for Ridge & Lasso
set.seed(16565)
detach(College)
attach(Boston)
bos.train <- sample_frac(Boston, 0.8)
```



```

bos.test = setdiff(Boston, bos.train)
nrow(bos.train) + nrow(bos.test) == nrow(Boston)

```

```
## [1] TRUE
```

```

rownames(bos.train) <- c()
rownames(bos.test) <- c()
mat.train <- model.matrix(crim ~ ., data = bos.train)[,-1]
mat.test <- model.matrix(crim ~ ., data=bos.test)[,-1]
summary(Boston)

```

```

##      crim      zn      indus      chas
##  Min.   : 0.00632   Min.   : 0.00   Min.   : 0.46   Min.   :0.00000
## 1st Qu.: 0.08204   1st Qu.: 0.00   1st Qu.: 5.19   1st Qu.:0.00000
## Median : 0.25651   Median : 0.00   Median : 9.69   Median :0.00000
## Mean   : 3.61352   Mean   : 11.36   Mean   :11.14   Mean   :0.06917
## 3rd Qu.: 3.67708   3rd Qu.: 12.50   3rd Qu.:18.10   3rd Qu.:0.00000
## Max.   :88.97620   Max.   :100.00   Max.   :27.74   Max.   :1.00000
##      nox      rm      age      dis
##  Min.   :0.3850   Min.   :3.561   Min.   : 2.90   Min.   : 1.130
## 1st Qu.:0.4490   1st Qu.:5.886   1st Qu.: 45.02   1st Qu.: 2.100
## Median :0.5380   Median :6.208   Median : 77.50   Median : 3.207
## Mean   :0.5547   Mean   :6.285   Mean   : 68.57   Mean   : 3.795
## 3rd Qu.:0.6240   3rd Qu.:6.623   3rd Qu.: 94.08   3rd Qu.: 5.188
## Max.   :0.8710   Max.   :8.780   Max.   :100.00   Max.   :12.127
##      rad      tax      ptratio      black
##  Min.   : 1.000   Min.   :187.0   Min.   :12.60   Min.   : 0.32
## 1st Qu.: 4.000   1st Qu.:279.0   1st Qu.:17.40   1st Qu.:375.38
## Median : 5.000   Median :330.0   Median :19.05   Median :391.44
## Mean   : 9.549   Mean   :408.2   Mean   :18.46   Mean   :356.67
## 3rd Qu.:24.000   3rd Qu.:666.0   3rd Qu.:20.20   3rd Qu.:396.23
## Max.   :24.000   Max.   :711.0   Max.   :22.00   Max.   :396.90
##      lstat      medv
##  Min.   : 1.73   Min.   : 5.00
## 1st Qu.: 6.95   1st Qu.:17.02
## Median :11.36   Median :21.20
## Mean   :12.65   Mean   :22.53
## 3rd Qu.:16.95   3rd Qu.:25.00
## Max.   :37.97   Max.   :50.00

```

```
str(Boston)
```

```

## 'data.frame': 506 obs. of 14 variables:
## $ crim : num 0.00632 0.02731 0.02729 0.03237 0.06905 ...
## $ zn : num 18 0 0 0 0 12.5 12.5 12.5 12.5 ...
## $ indus : num 2.31 7.07 7.07 2.18 2.18 2.18 7.87 7.87 7.87 7.87 ...
## $ chas : int 0 0 0 0 0 0 0 0 0 0 ...
## $ nox : num 0.538 0.469 0.469 0.458 0.458 0.458 0.524 0.524 0.524 0.524 ...
## $ rm : num 6.58 6.42 7.18 7 7.15 ...
## $ age : num 65.2 78.9 61.1 45.8 54.2 58.7 66.6 96.1 100 85.9 ...
## $ dis : num 4.09 4.97 4.97 6.06 6.06 ...
## $ rad : int 1 2 2 3 3 3 5 5 5 5 ...
## $ tax : num 296 242 242 222 222 222 311 311 311 311 ...
## $ ptratio: num 15.3 17.8 17.8 18.7 18.7 18.7 15.2 15.2 15.2 15.2 ...
## $ black : num 397 397 393 395 397 ...

```

```
## $ lstat : num 4.98 9.14 4.03 2.94 5.33 ...
## $ medv : num 24 21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 ...
```

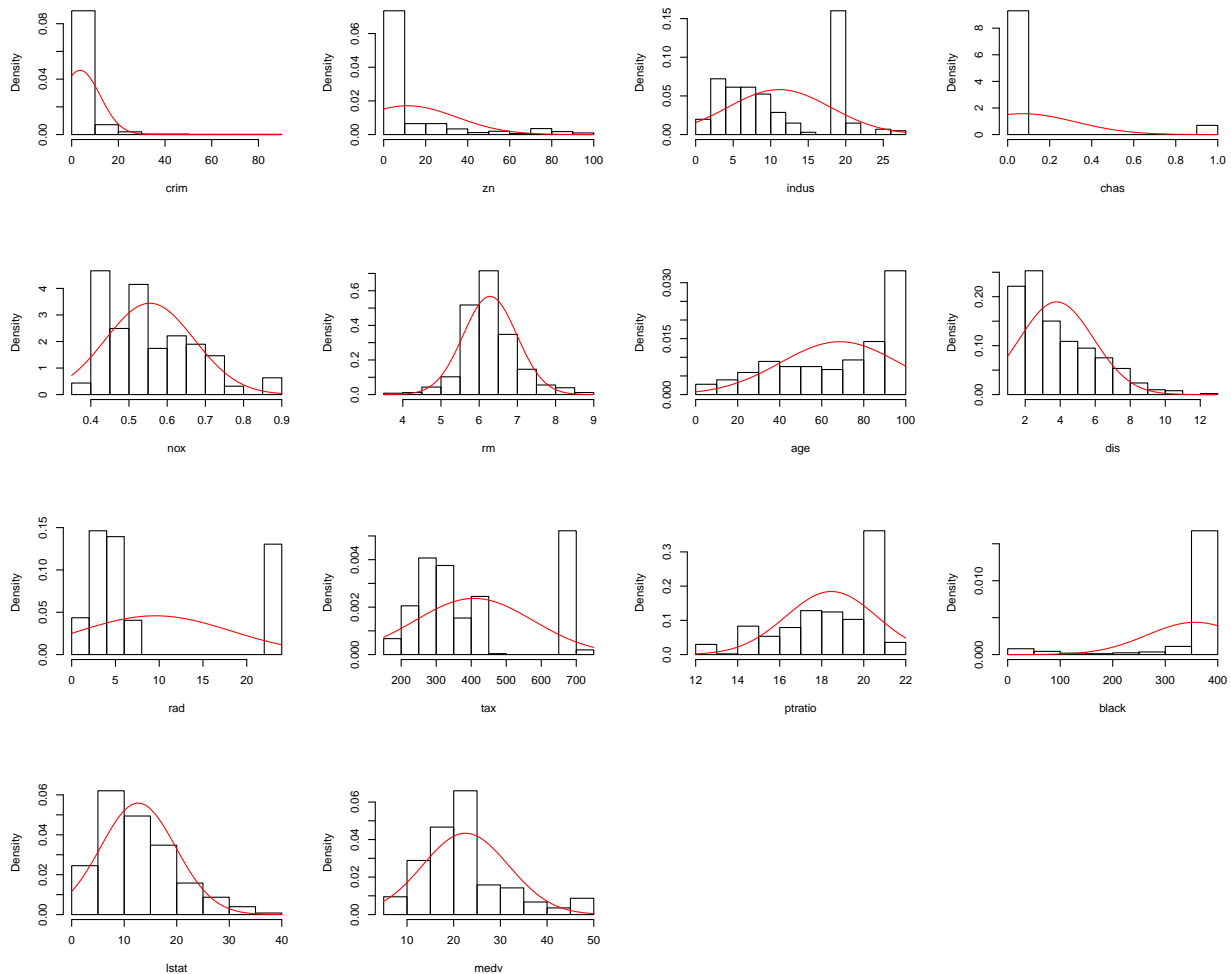
```
anyDuplicated(Boston)
```

```
## [1] 0
```


```
sum(is.na(Boston))
```

```
## [1] 0
```

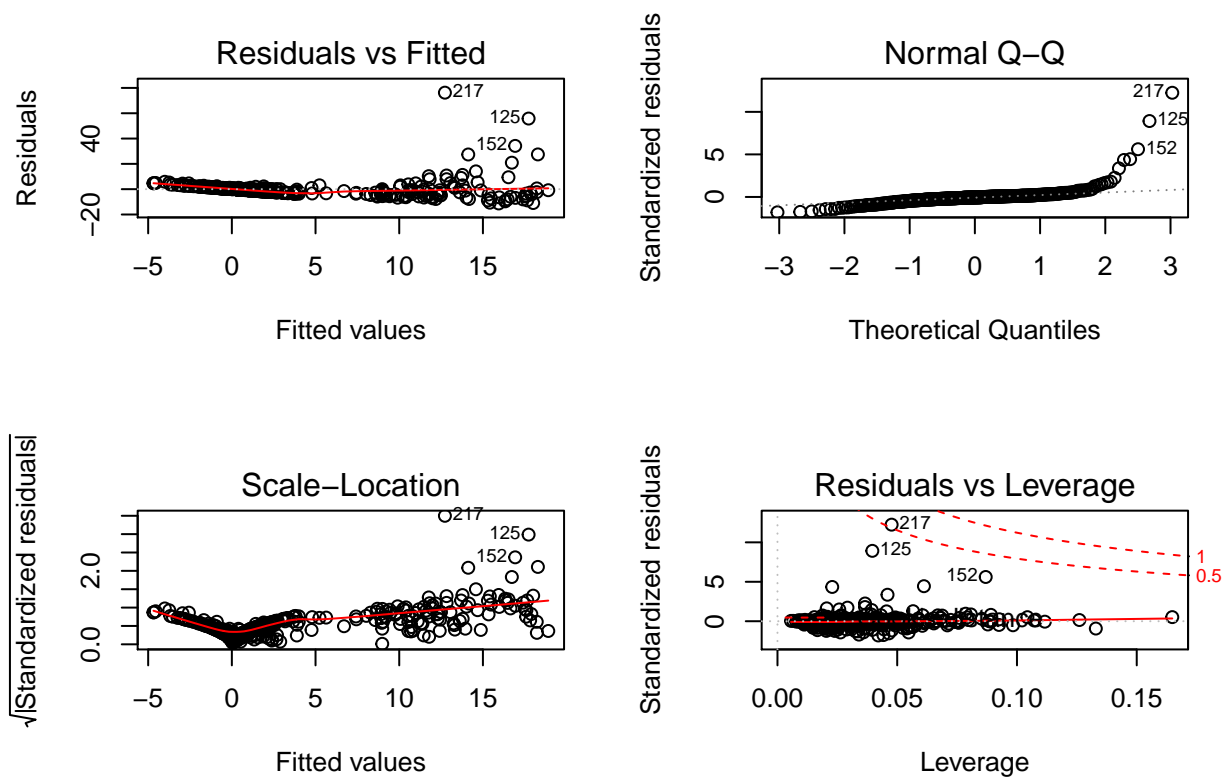
```
uniPlot(Boston, type = "histogram")
```



```
r <- cor(Boston)
title <- 'Boston Data'
corrplot::corrplot(r, method = "color", type = 'upper', diag = FALSE, addCoef.col = "black",
  order = "hclust", title = title, mar=c(0,0,1,0))
```




```
lm.fit <- lm(crim ~ ., data = bos.train)
par(mfrow=c(2,2))
plot(lm.fit)
```



```
lm.pred <- predict(lm.fit, bos.test)
lm.er <- mean((bos.test$crim - lm.pred)^2); lm.er
```

```
## [1] 45.90214
```

```
grid = 10^seq(10, -2, length = 100)
ridge.fit <- cv.glmnet(mat.train, bos.train$crim, alpha = 0, lambda = grid, thresh = 1e-12)
plot(ridge.fit)
```



```
ridge.fit$lambda.min
```

```
## [1] 0.6579332
```

```
ridge.pred <- predict(ridge.fit, newx = mat.test, s = ridge.fit$lambda.min)
```

```
ridge.er <- mean((bos.test$crim - ridge.pred)^2); ridge.er
```


```
## [1] 47.24654
```

```
predict(ridge.fit, s = ridge.fit$lambda.min, type = "coefficients")
```

```
## 14 x 1 sparse Matrix of class "dgCMatrix"
```

```
##              1
## (Intercept)  9.674512475
## zn          0.026115285
## indus       -0.061593845
## chas        -0.874509207
## nox         -4.886593189
## rm          0.211496504
## age         0.002191499
## dis        -0.633453558
## rad         0.387455204
## tax         0.003302351
## ptratio    -0.066763814
## black      -0.014351994
## lstat       0.117045867
## medv       -0.096758718
```

```
lass.fit = cv.glmnet(mat.train, bos.train$crim, alpha = 1, lambda = grid, thresh = 1e-12)
plot(lass.fit)
```




```
lass.pred = predict(lass.fit, s = lass.fit$lambda.min, newx = mat.test)
lass.er = mean((bos.test$crim - lass.pred)^2)
predict(lass.fit, s = lass.fit$lambda.min, type = "coefficients")
```

```
## 14 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept) 10.25170685
## zn          0.02477619
## indus       -0.03161461
## chas        -0.71055271
## nox         -4.10731848
## rm          .
## age         .
## dis        -0.56965751
## rad         0.46463310
## tax         .
## ptratio    -0.04737249
## black      -0.01407236
## lstat       0.09664408
## medv       -0.08298947
```

```
library(pls)
pcr.fit <- pcr(crim ~ ., data = bos.train, scale = TRUE, validation = "CV")
validationplot(pcr.mod, val.type = 'MSEP')
```

Apps



```
summary(pcr.mod)
```

```
## Data:      X dimension: 622 17
## Y dimension: 622 1
## Fit method: svdpc
## Number of components considered: 17
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              3625    3485    1720    1729    1373    1323    1294
## adjCV           3625    3484    1718    1730    1327    1307    1291
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV          1264    1213    1199    1198    1199    1200    1204
## adjCV        1263    1209    1196    1196    1197    1197    1201
##      14 comps 15 comps 16 comps 17 comps
## CV          1204    1206    1028    1019
## adjCV        1201    1203    1025    1016
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps
## X          31.540  57.13   64.18   69.91   75.36   80.53   84.52
## Apps       7.907  77.82   77.90   87.08   87.61   87.77   88.38
##      8 comps  9 comps 10 comps 11 comps 12 comps 13 comps 14 comps
## X          88.04  90.96   93.17   95.27   97.07   98.16   98.94
```

```
## Apps      89.24      89.70      89.74      89.82      89.86      89.86      89.88
##          15 comps  16 comps  17 comps
## X          99.45      99.86     100.00
## Apps      90.17      92.60      92.92
```


```
pcr.pred <- predict(pcr.fit, bos.test, ncomp = 10)
pcr.er <- mean((bos.test$crim - pcr.pred)^2); pcr.er
```

```
## [1] 49.54782
```

```
pls.fit <- plsr(crim ~ ., data = bos.train, scale = TRUE, validation = "CV")
summary(pls.fit)
```

```
## Data:      X dimension: 405 13
## Y dimension: 405 1
## Fit method: kernelppls
## Number of components considered: 13
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              8.521    6.951    6.573    6.573    6.551    6.541    6.540
## adjCV           8.521    6.948    6.567    6.557    6.539    6.529    6.525
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV          6.560    6.534    6.530    6.529    6.528    6.528    6.528
## adjCV       6.542    6.520    6.516    6.515    6.514    6.515    6.515
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps
## X          47.89    57.48    61.20    71.80    77.24    80.44    82.82
## crim       34.64    42.65    44.53    44.87    45.18    45.47    45.55
##      8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## X          86.54    90.25    94.84    97.26    98.51    100.00
## crim       45.57    45.57    45.57    45.57    45.57    45.57
```

```
validationplot(pls.fit, val.type = 'MSEP')
```





```
pls.pred <- predict(pls.fit, bos.test, ncomp = 9)
pls.er <- mean((bos.test$crim - pls.pred)^2); pls.er
```


```
## [1] 45.82656
```

(B)

```
mod.errors <- c(lm.er, ridge.er, lass.er, pcr.er, pls.er)
names(mod.errors) <- c('OLS', 'Ridge', 'Lasso', 'PCR', 'PLS')
barplot(mod.errors, main = 'Test Error by Method', xlab = 'Method', ylab = 'RSS')
```



```
plot(mod.errors, type = 'h', col='blue', xaxt='n')  
axis(1, at=1:5, lab=c('OLS', 'Ridge', 'Lasso', 'PCR', 'PLS'))
```




```
kable(mod.errors, col.names = 'RSS')%>%
  kable_styling()
```

| | RSS |
|-------|----------|
| OLS | 45.90214 |
| Ridge | 47.24654 |
| Lasso | 47.33744 |
| PCR | 49.54782 |
| PLS | 45.82656 |

#Again, calculating all the r2 in this manner was inspired by a forum post

```
t.avg <- mean(bos.test$crim)
ols.r2 = 1 - mean((lm.pred - bos.test$crim)^2) / mean((t.avg - bos.test$crim)^2)
ridge.r2 = 1 - mean((ridge.pred - bos.test$crim)^2) / mean((t.avg - bos.test$crim)^2)
lass.r2 = 1 - mean((lass.pred - bos.test$crim)^2) / mean((t.avg - bos.test$crim)^2)
pcr.r2 = 1 - mean((pcr.pred - bos.test$crim)^2) / mean((t.avg - bos.test$crim)^2)
pls.r2 = 1 - mean((pls.pred - bos.test$crim)^2) / mean((t.avg - bos.test$crim)^2)

mod.effect <- c(ols.r2, ridge.r2, lass.r2, pcr.r2, pls.r2)
names(mod.effect) <- c('OLS', 'Ridge', 'Lasso', 'PCR', 'PLS')
plot(mod.effect, type = 'h', col='blue', xaxt='n')
axis(1, at=1:5, lab=c('OLS', 'Ridge', 'Lasso', 'PCR', 'PLS'))
```



```
kable(mod.effect, col.names = 'r2')%>%
  kable_styling()
```

| | r2 |
|-------|-----------|
| OLS | 0.4272306 |
| Ridge | 0.4104551 |
| Lasso | 0.4093208 |
| PCR | 0.3817395 |
| PLS | 0.4281737 |

OLS was used as a basis for comparison. All of the models perform fairly similarly, save PCR which stood out as the worst model (i.e. largest test error, lowest r^2). PLS had the smallest test error, though not smaller than OLS which had marginally better test error and r^2 . Of the methods from chapter 6 (excepting step wise and subset selection), the PLS regression model performs the best, accounting for 39% of the variation in test set crime.

(C)

The PLS model that performed best predicting Boston's crime fit 9 linear combinations of variables; while this reduced dimensions, it nonetheless utilized much of the information available. Thus, I would not remove any features.