

PQHS 471 HW 3

Gregory Powers

March 18, 2018

Chapter 7.9

(a)

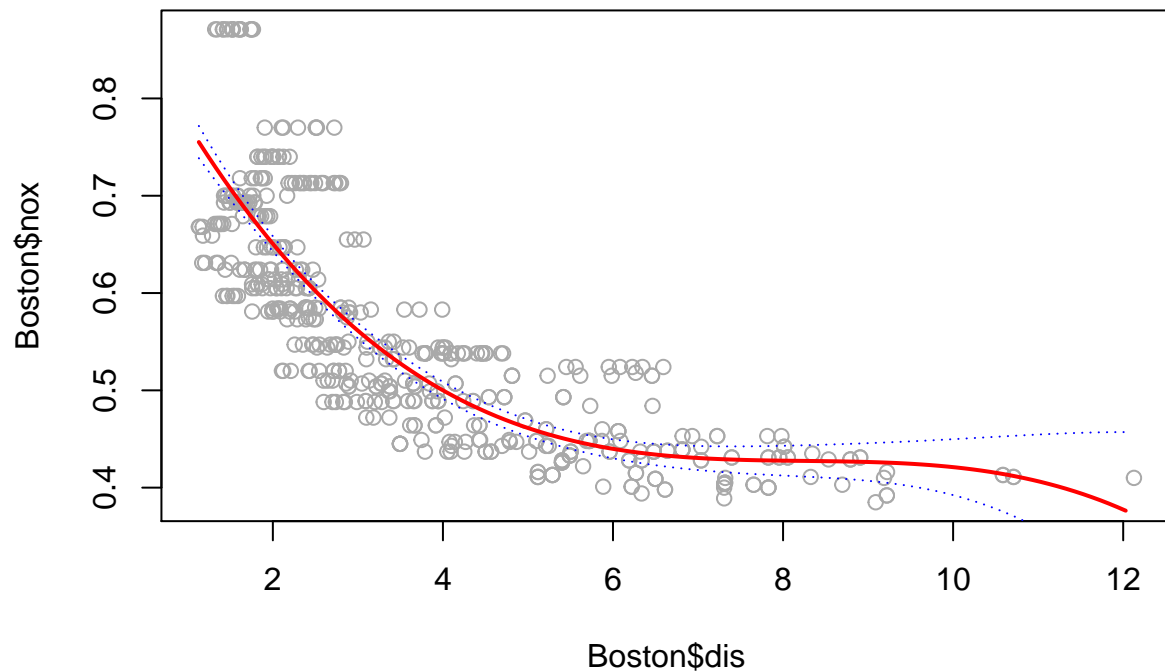
```
library(ISLR)
library(MASS)
library(ggplot2)

lm.1 <- lm(nox~poly(dis, 3), data = Boston)
summary(lm.1)

##
## Call:
## lm(formula = nox ~ poly(dis, 3), data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.121130 -0.040619 -0.009738  0.023385  0.194904
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   0.554695   0.002759  201.021 < 2e-16 ***
## poly(dis, 3)1 -2.003096   0.062071  -32.271 < 2e-16 ***
## poly(dis, 3)2  0.856330   0.062071   13.796 < 2e-16 ***
## poly(dis, 3)3 -0.318049   0.062071   -5.124 4.27e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.06207 on 502 degrees of freedom
## Multiple R-squared:  0.7148, Adjusted R-squared:  0.7131
## F-statistic: 419.3 on 3 and 502 DF,  p-value: < 2.2e-16

dis.lim <- range(Boston$dis)
dis.grid <- seq(from = dis.lim[1], to = dis.lim[2], by = .1)
pred <- predict(lm.1, newdata = list(dis=dis.grid), se = TRUE)
se.bands <- cbind(pred$fit + 2*pred$se.fit, pred$fit - 2*pred$se.fit)
plot(Boston$nox~Boston$dis, xlim = dis.lim, cex = 1, col = "darkgrey")
title("Nox~Dis Cubic Fit")
lines(dis.grid, pred$fit, lwd = 2, col = "red")
matlines(dis.grid, se.bands, lwd = 1, col = "blue", lty = 3)
```

Nox~Dis Cubic Fit




The cubic fit of $\text{nox} \sim \text{dis}$ is good: all terms are significant and together account for 71% of the variance in nox . However, we can see from the plot that while the function fits well throughout most of the range of dis , the confidence bands grow larger as dis approaches 12, suggesting the presence of outliers.

(b)

```
rss = rep(0, 10)
for (i in 1:10) {
  lm.fit <- lm(nox ~ poly(dis, i), data = Boston)
  rss[i] <- sum(lm.fit$residuals^2)
}
rss
```


```
## [1] 2.768563 2.035262 1.934107 1.932981 1.915290 1.878257 1.849484
## [8] 1.835630 1.833331 1.832171
```

```
plot(rss, type = "b", xlab = "Degree", lwd = 2)
```



```
library(boot)
set.seed(1)
poly.cv <- rep(0, 10)
for (i in 1:10) {
  glm.fit <- glm(nox ~ poly(dis, i), data = Boston)
  poly.cv[i] <- cv.glm(Boston, glm.fit, K = 10)$delta[1] # note to self: [1] = standard, [2] = bias-
}

plot(poly.cv, type = "b", lwd = 2, ylab = "MSE", xlab = "Poly Degree"); poly.cv
```



```
## [1] 0.005536329 0.004077147 0.003899587 0.003862127 0.004298590
## [6] 0.005095283 0.013680327 0.005284520 0.013355413 0.004148996
```

Judging from the graph and the poly.cv output, a 4th degree polynomial fit minimizes training error, albeit by a tiny margin over a 3rd degree fit.

```
library(splines)
#length(nox)
fit.sp <- lm(nox~bs(dis, df = 4, knots = c(2,3,5)), data = Boston)
summary(fit.sp)
```

```
##
## Call:
## lm(formula = nox ~ bs(dis, df = 4, knots = c(2, 3, 5)), data = Boston)
##
## Residuals:
```

| | Min | 1Q | Median | 3Q | Max |
|--|----------|----------|----------|---------|---------|
| | -0.12778 | -0.03834 | -0.01019 | 0.02273 | 0.19509 |


```
##
## Coefficients:
```

| | Estimate | Std. Error | t value | Pr(> t) |
|---|-----------|------------|---------|----------|
| ## (Intercept) | 0.649937 | 0.024948 | 26.052 | < 2e-16 |
| ## bs(dis, df = 4, knots = c(2, 3, 5))1 | 0.107267 | 0.036455 | 2.942 | 0.00341 |
| ## bs(dis, df = 4, knots = c(2, 3, 5))2 | -0.003582 | 0.024669 | -0.145 | 0.88461 |
| ## bs(dis, df = 4, knots = c(2, 3, 5))3 | -0.146759 | 0.028851 | -5.087 | 5.16e-07 |
| ## bs(dis, df = 4, knots = c(2, 3, 5))4 | -0.222917 | 0.034193 | -6.519 | 1.73e-10 |
| ## bs(dis, df = 4, knots = c(2, 3, 5))5 | -0.254819 | 0.049313 | -5.167 | 3.44e-07 |

```
## bs(dis, df = 4, knots = c(2, 3, 5))6 -0.234438 0.054364 -4.312 1.95e-05
##
## (Intercept) ***
## bs(dis, df = 4, knots = c(2, 3, 5))1 **
## bs(dis, df = 4, knots = c(2, 3, 5))2
## bs(dis, df = 4, knots = c(2, 3, 5))3 ***
## bs(dis, df = 4, knots = c(2, 3, 5))4 ***
## bs(dis, df = 4, knots = c(2, 3, 5))5 ***
## bs(dis, df = 4, knots = c(2, 3, 5))6 ***
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.06057 on 499 degrees of freedom
## Multiple R-squared: 0.7301, Adjusted R-squared: 0.7268
## F-statistic: 224.9 on 6 and 499 DF, p-value: < 2.2e-16

pred.sp <- predict(fit.sp, newdata = list(dis=dis.grid), se = TRUE)
se.bands.sp <- cbind(pred.sp$fit + 2*pred.sp$se.fit, pred.sp$fit - 2*pred.sp$se.fit)
plot(Boston$nox~Boston$dis, cex = 1, col = "darkgrey")
title("Nox~Dis 4 DF Spline")
lines(dis.grid, pred.sp$fit, lwd = 2, col = "red")
matlines(dis.grid, se.bands.sp, lwd = 1, col = "blue", lty = 3)
```

Nox~Dis 4 DF Spline



Knots were selected by summary statistics: the first, second and third quantiles.

(e)


```

rss.sp <- rep(NA, 18)
for (i in 3:18) {
  fit.sp.rss <- lm(nox~bs(dis, df = i), data = Boston)
  rss.sp[i] <- sum(fit.sp.rss$residuals^2)
}
rss.sp <- rss.sp[-c(1,2)]; rss.sp

## [1] 1.934107 1.922775 1.840173 1.833966 1.829884 1.816995 1.825653
## [8] 1.792535 1.796992 1.788999 1.782350 1.781838 1.782798 1.783546
## [15] 1.779789 1.775838

x <- c(3:18)
plot(3:18, rss.sp, type = "b", ylab = "RSS", xlab = "# DF", xaxt='n')
axis(1, at = x)

```



Error decreases more or less as flexibility increases, though not monotonically. There is a small increase DF 8 - DF 9, and again 10-11.

(f)

```

set.seed(1)
bs.cv <- rep(NA, 18)
for (i in 3:18) {
  bs.fit <- glm(nox ~ bs(dis, df = i), data = Boston)
  bs.cv[i] <- cv.glm(Boston, bs.fit, K = 10)$delta[1] # note to self: [1] = standard, [2] = bias-cor
}
bs.cv <- bs.cv[-c(1,2)]; bs.cv

```

```
## [1] 0.003865662 0.003915074 0.003732840 0.003682221 0.003726081
## [6] 0.003701344 0.003741995 0.003701813 0.003720982 0.003654150
## [11] 0.003701321 0.003768401 0.003739480 0.003739913 0.003804258
## [16] 0.003775504
```

```
plot(bs.cv, type = "b", lwd = 2, ylab = "Error", xlab = "Spline DF", xaxt='n')
x <- c(1:18)
axis(1, at = x)
```



The minimum CV error is where $DF = 10$, though this varies depending on the RNG and may not be true if the set seed value is something other than 1. `cv.glm` issues a number of warnings which are suppressed for a more tidy output.

Chapter 8.9

(a)

```
summary(OJ)
```

| ## | Purchase | WeekofPurchase | StoreID | PriceCH | PriceMM |
|----|----------|----------------|--------------|---------------|---------------|
| ## | CH:653 | Min. :227.0 | Min. :1.00 | Min. :1.690 | Min. :1.690 |
| ## | MM:417 | 1st Qu.:240.0 | 1st Qu.:2.00 | 1st Qu.:1.790 | 1st Qu.:1.990 |
| ## | | Median :257.0 | Median :3.00 | Median :1.860 | Median :2.090 |
| ## | | Mean :254.4 | Mean :3.96 | Mean :1.867 | Mean :2.085 |
| ## | | 3rd Qu.:268.0 | 3rd Qu.:7.00 | 3rd Qu.:1.990 | 3rd Qu.:2.180 |
| ## | | Max. :278.0 | Max. :7.00 | Max. :2.090 | Max. :2.290 |

```
##      DiscCH      DiscMM      SpecialCH      SpecialMM
## Min.   :0.00000 Min.   :0.0000 Min.   :0.0000 Min.   :0.0000
## 1st Qu.:0.00000 1st Qu.:0.0000 1st Qu.:0.0000 1st Qu.:0.0000
## Median :0.00000 Median :0.0000 Median :0.0000 Median :0.0000
## Mean   :0.05186 Mean   :0.1234 Mean   :0.1477 Mean   :0.1617
## 3rd Qu.:0.00000 3rd Qu.:0.2300 3rd Qu.:0.0000 3rd Qu.:0.0000
## Max.   :0.50000 Max.   :0.8000 Max.   :1.0000 Max.   :1.0000
##      LoyalCH      SalePriceMM      SalePriceCH      PriceDiff
## Min.   :0.000011 Min.   :1.190 Min.   :1.390 Min.   : -0.6700
## 1st Qu.:0.325257 1st Qu.:1.690 1st Qu.:1.750 1st Qu.: 0.0000
## Median :0.600000 Median :2.090 Median :1.860 Median : 0.2300
## Mean   :0.565782 Mean   :1.962 Mean   :1.816 Mean   : 0.1465
## 3rd Qu.:0.850873 3rd Qu.:2.130 3rd Qu.:1.890 3rd Qu.: 0.3200
## Max.   :0.999947 Max.   :2.290 Max.   :2.090 Max.   : 0.6400
##      Store7      PctDiscMM      PctDiscCH      ListPriceDiff
## No :714 Min.   :0.0000 Min.   :0.00000 Min.   :0.000
## Yes:356 1st Qu.:0.0000 1st Qu.:0.00000 1st Qu.:0.140
##      Median :0.0000 Median :0.00000 Median :0.240
##      Mean   :0.0593 Mean   :0.02731 Mean   :0.218
##      3rd Qu.:0.1127 3rd Qu.:0.00000 3rd Qu.:0.300
##      Max.   :0.4020 Max.   :0.25269 Max.   :0.440
##      STORE
## Min.   :0.000
## 1st Qu.:0.000
## Median :2.000
## Mean   :1.631
## 3rd Qu.:3.000
## Max.   :4.000
```

```
str(OJ)
```

```
## 'data.frame': 1070 obs. of 18 variables:
## $ Purchase : Factor w/ 2 levels "CH","MM": 1 1 1 2 1 1 1 1 1 1 ...
## $ WeekofPurchase: num 237 239 245 227 228 230 232 234 235 238 ...
## $ StoreID : num 1 1 1 1 7 7 7 7 7 7 ...
## $ PriceCH : num 1.75 1.75 1.86 1.69 1.69 1.69 1.69 1.75 1.75 1.75 ...
## $ PriceMM : num 1.99 1.99 2.09 1.69 1.69 1.99 1.99 1.99 1.99 1.99 ...
## $ DiscCH : num 0 0 0.17 0 0 0 0 0 0 0 ...
## $ DiscMM : num 0 0.3 0 0 0 0 0.4 0.4 0.4 0.4 ...
## $ SpecialCH : num 0 0 0 0 0 0 1 1 0 0 ...
## $ SpecialMM : num 0 1 0 0 0 1 1 0 0 0 ...
## $ LoyalCH : num 0.5 0.6 0.68 0.4 0.957 ...
## $ SalePriceMM : num 1.99 1.69 2.09 1.69 1.69 1.99 1.59 1.59 1.59 1.59 ...
## $ SalePriceCH : num 1.75 1.75 1.69 1.69 1.69 1.69 1.69 1.75 1.75 1.75 ...
## $ PriceDiff : num 0.24 -0.06 0.4 0 0 0.3 -0.1 -0.16 -0.16 -0.16 ...
## $ Store7 : Factor w/ 2 levels "No","Yes": 1 1 1 1 2 2 2 2 2 2 ...
## $ PctDiscMM : num 0 0.151 0 0 0 ...
## $ PctDiscCH : num 0 0 0.0914 0 0 ...
## $ ListPriceDiff : num 0.24 0.24 0.23 0 0 0.3 0.3 0.24 0.24 0.24 ...
## $ STORE : num 1 1 1 1 0 0 0 0 0 0 ...
```

```
anyDuplicated(OJ)
```

```
## [1] 439
```



```
sum(is.na(OJ))
```

```
## [1] 0
```

```
set.seed(1)
train <- sample(1:nrow(OJ), 800)
oj.train <- OJ[train,]
oj.test <- OJ[-train,]
```

(b)

```
library(tree)
tree.fit <- tree(Purchase ~ ., data = oj.train)
summary(tree.fit)
```

```
##
## Classification tree:
## tree(formula = Purchase ~ ., data = oj.train)
## Variables actually used in tree construction:
## [1] "LoyalCH"      "PriceDiff"    "SpecialCH"    "ListPriceDiff"
## Number of terminal nodes: 8
## Residual mean deviance: 0.7305 = 578.6 / 792
## Misclassification error rate: 0.165 = 132 / 800
```

The tree has 8 nodes and a training error rate of 0.165

(c)


```
tree.fit

## node), split, n, deviance, yval, (yprob)
##      * denotes terminal node
##
## 1) root 800 1064.00 CH ( 0.61750 0.38250 )
##    2) LoyalCH < 0.508643 350 409.30 MM ( 0.27143 0.72857 )
##      4) LoyalCH < 0.264232 166 122.10 MM ( 0.12048 0.87952 )
##        8) LoyalCH < 0.0356415 57 10.07 MM ( 0.01754 0.98246 ) *
##        9) LoyalCH > 0.0356415 109 100.90 MM ( 0.17431 0.82569 ) *
##      5) LoyalCH > 0.264232 184 248.80 MM ( 0.40761 0.59239 )
##        10) PriceDiff < 0.195 83 91.66 MM ( 0.24096 0.75904 )
##          20) SpecialCH < 0.5 70 60.89 MM ( 0.15714 0.84286 ) *
##          21) SpecialCH > 0.5 13 16.05 CH ( 0.69231 0.30769 ) *
##        11) PriceDiff > 0.195 101 139.20 CH ( 0.54455 0.45545 ) *
##    3) LoyalCH > 0.508643 450 318.10 CH ( 0.88667 0.11333 )
##      6) LoyalCH < 0.764572 172 188.90 CH ( 0.76163 0.23837 )
##        12) ListPriceDiff < 0.235 70 95.61 CH ( 0.57143 0.42857 ) *
##        13) ListPriceDiff > 0.235 102 69.76 CH ( 0.89216 0.10784 ) *
##      7) LoyalCH > 0.764572 278 86.14 CH ( 0.96403 0.03597 ) *
```

20 is a terminal node that predicts MM. The split criterion is $\text{SpecialCH} < .5$. The branch has 70 observations with a deviance of 60.89. About 16% of the branch has Sales = CH, 84% Sales = MM.

(d)

```
plot(tree.fit)
text(tree.fit)
```



LoyalCH appears to be the most important predictor of Purchase.

(e)

```
tree.pred <- predict(tree.fit, oj.test, type = "class")
table(tree.pred, oj.test$Purchase)
```

```
##
## tree.pred CH MM
##          CH 147 49
##          MM  12 62
```

```
1-(209)/270
```

```
## [1] 0.2259259
```

Our test error rate is about 23%.

(f)


```
cv.oj <- cv.tree(tree.fit, FUN = prune.misclass); cv.oj
```

```
## $size
## [1] 8 5 2 1
##
## $dev
## [1] 156 156 156 306
##
## $k
## [1] -Inf 0.000000 4.666667 160.000000
```

```
##
## $method
## [1] "misclass"
##
## attr("class")
## [1] "prune"          "tree.sequence"
```

(g)


```
plot(cv.oj$size, cv.oj$dev, type = "b", lwd = 2, ylab = "Error", xlab = "Tree Size")
```



(h) A tree size of 5 is associated with the lowest error rate.

(i)

```
prune.fit <- prune.misclass(tree.fit, best = 5)
plot(prune.fit)
text(prune.fit)
```



(j)

```
summary(tree.fit)
```

```
##
## Classification tree:
## tree(formula = Purchase ~ ., data = oj.train)
## Variables actually used in tree construction:
## [1] "LoyalCH"      "PriceDiff"    "SpecialCH"    "ListPriceDiff"
## Number of terminal nodes: 8
## Residual mean deviance: 0.7305 = 578.6 / 792
## Misclassification error rate: 0.165 = 132 / 800
```

```
summary(prune.fit)
```

```
##
## Classification tree:
## snip.tree(tree = tree.fit, nodes = 3:4)
## Variables actually used in tree construction:
## [1] "LoyalCH"      "PriceDiff"    "SpecialCH"
## Number of terminal nodes: 5
## Residual mean deviance: 0.8256 = 656.4 / 795
## Misclassification error rate: 0.165 = 132 / 800
```

The train error rates are identical.

(k)

```
prune.pred <- predict(prune.fit, oj.test, type = "class")
table(prune.pred, oj.test$Purchase)
```

```
##
## prune.pred  CH  MM
##           CH 147  49
##           MM  12  62
```

```
1-(209)/270
```

```
## [1] 0.2259259
```

The test error rate of the pruned and unpruned models are identical.

9.1.6 Khan Data

```
summary(Khan)
```

```
##           Length Class  Mode
## xtrain 145404 -none- numeric
## xtest   46160 -none- numeric
## ytrain    63 -none- numeric
## ytest    20 -none- numeric
```

```
str(Khan)
```

```
## List of 4
## $ xtrain: num [1:63, 1:2308] 0.7733 -0.0782 -0.0845 0.9656 0.0757 ...
##   ..- attr(*, "dimnames")=List of 2
##     .. ..$ : chr [1:63] "V1" "V2" "V3" "V4" ...
##     .. ..$ : NULL
## $ xtest : num [1:20, 1:2308] 0.14 1.164 0.841 0.685 -1.956 ...
##   ..- attr(*, "dimnames")=List of 2
##     .. ..$ : chr [1:20] "V1" "V2" "V4" "V6" ...
##     .. ..$ : NULL
## $ ytrain: num [1:63] 2 2 2 2 2 2 2 2 2 2 ...
## $ ytest : num [1:20] 3 2 4 2 1 3 4 2 3 1 ...
```

```
anyDuplicated(Khan)
```

```
## [1] 0
```

```
sum(is.na(Khan))
```

```
## [1] 0
```

```
table(Khan$ytrain)
```

```
##
##  1  2  3  4
##  8 23 12 20
```

```
khan.train <- data.frame(x = Khan$xtrain, y = as.factor(Khan$ytrain))
dim(khan.train)
```

```
## [1]    63 2309
```

```

khan.test <- data.frame(x = Khan$xtest, y = as.factor(Khan$ytest))
dim(khan.test)

## [1] 20 2309

library(parallel)
library(doParallel)

## Loading required package: foreach
## Loading required package: iterators
library(caret)

## Loading required package: lattice

##
## Attaching package: 'lattice'

## The following object is masked from 'package:boot':
##
##      melanoma

set.seed(1)
cluster <- makeCluster(detectCores() - 1)
registerDoParallel(cluster)
tunegrid <- expand.grid(mtry=c(1:10))
train.control <- trainControl(method="cv", number=10, search = "grid", allowParallel = TRUE)
model.for <- train(y~., data=khan.train, method="rf", tuneLength = 15, truneGrid = tunegrid, trControl=
stopCluster(cluster)
registerDoSEQ()

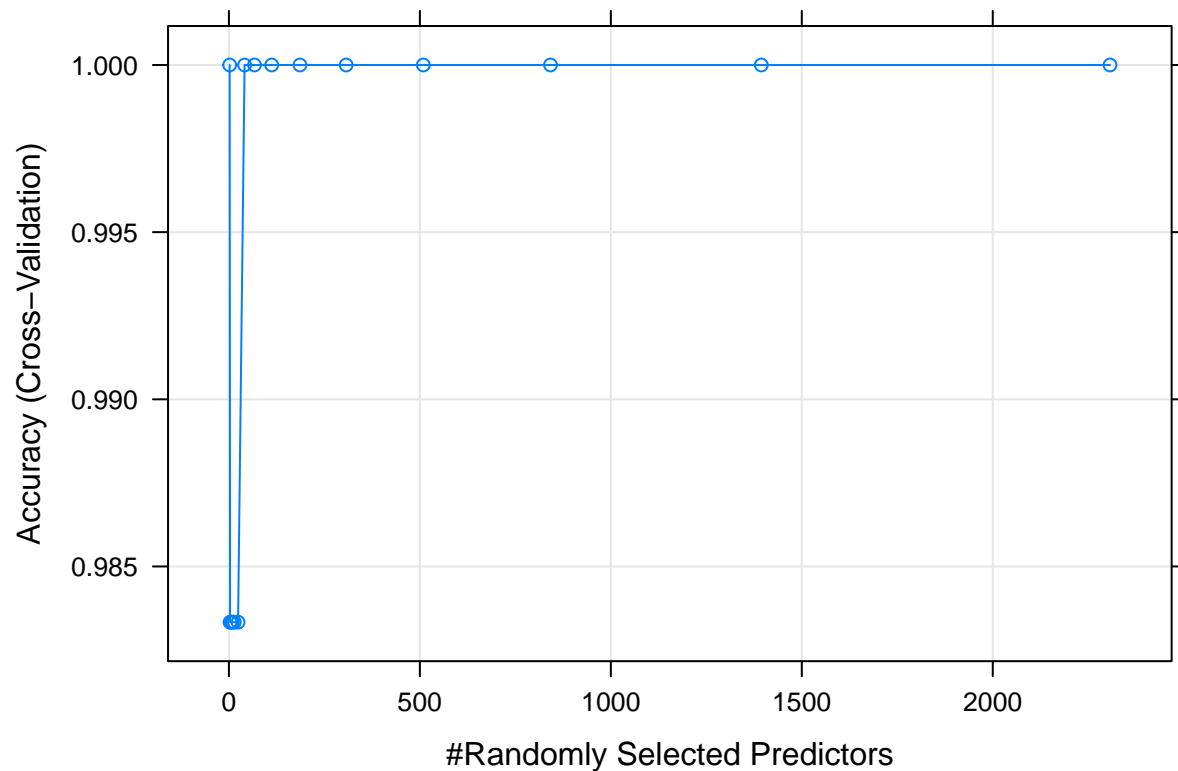
print(model.for)

## Random Forest
##
## 63 samples
## 2308 predictors
## 4 classes: '1', '2', '3', '4'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 57, 56, 57, 56, 57, 58, ...
## Resampling results across tuning parameters:
##
##  mtry  Accuracy  Kappa
##    2  1.0000000  1.0000000
##    3  0.9833333  0.9769231
##    5  0.9833333  0.9769231
##    9  0.9833333  0.9769231
##   14  0.9833333  0.9769231
##   24  0.9833333  0.9769231
##   41  1.0000000  1.0000000
##   67  1.0000000  1.0000000
##  112  1.0000000  1.0000000
##  186  1.0000000  1.0000000
##  307  1.0000000  1.0000000
##  509  1.0000000  1.0000000

```

```
##      842  1.0000000  1.0000000
##     1394  1.0000000  1.0000000
##     2307  1.0000000  1.0000000
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 2.
```

```
plot(model.for)
```



```
summary(model.for)
```

```
##          Length Class      Mode
## call           6  -none-    call
## type           1  -none- character
## predicted       63  factor   numeric
## err.rate      7500 -none-    numeric
## confusion       20 -none-    numeric
## votes         252 matrix    numeric
## oob.times       63 -none-    numeric
## classes         4  -none- character
## importance     2308 -none-    numeric
## importanceSD     0  -none-    NULL
## localImportance  0  -none-    NULL
## proximity       0  -none-    NULL
## ntree           1  -none-    numeric
## mtry            1  -none-    numeric
## forest         14 -none-    list
```

```
## y          63  factor    numeric
## test       0  -none-    NULL
## inbag      0  -none-    NULL
## xNames    2308 -none-    character
## problemType 1  -none-    character
## tuneValue  1  data.frame list
## obsLevels  4  -none-    character
## param      2  -none-    list

predict.for <- predict.train(object=model.for, khan.test, type="raw")
confusionMatrix(predict.for, khan.test$y)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction 1 2 3 4
##           1 3 0 0 0
##           2 0 5 1 0
##           3 0 0 1 0
##           4 0 1 4 5
##
## Overall Statistics
##
##           Accuracy : 0.7
##           95% CI : (0.4572, 0.8811)
##           No Information Rate : 0.3
##           P-Value [Acc > NIR] : 0.000261
##
##           Kappa : 0.5987
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: 1 Class: 2 Class: 3 Class: 4
## Sensitivity          1.00  0.8333  0.1667  1.0000
## Specificity          1.00  0.9286  1.0000  0.6667
## Pos Pred Value       1.00  0.8333  1.0000  0.5000
## Neg Pred Value       1.00  0.9286  0.7368  1.0000
## Prevalence           0.15  0.3000  0.3000  0.2500
## Detection Rate       0.15  0.2500  0.0500  0.2500
## Detection Prevalence 0.15  0.3000  0.0500  0.5000
## Balanced Accuracy    1.00  0.8810  0.5833  0.8333

varImp(model.for)
```


```
## rf variable importance
##
##   only 20 most important variables shown (out of 2308)
##
##           Overall
## x.545    100.00
## x.187     99.46
## x.607     93.78
## x.2050    92.42
## x.2046    85.50
```




```
## x.380      84.83
## x.417      79.73
## x.1389     79.12
## x.430      76.24
## x.1706     72.98
## x.1110     72.36
## x.554      71.26
## x.1055     71.22
## x.1645     71.10
## x.246      70.64
## x.335      70.37
## x.1159     70.30
## x.1194     69.65
## x.1105     68.44
## x.166      68.30
```

```
cluster <- makeCluster(detectCores() - 1)
registerDoParallel(cluster)
set.seed(1)
grid.xg = expand.grid(
  nrounds = 10,
  max_depth = c(3, 6, 10),
  eta = c(0.1, 0.01),
  gamma = c(0, 1, 2),
  colsample_bytree = c(0.4, 0.7, 1.0),
  min_child_weight = c(1, 1.5, 2),
  subsample = .5
)
#train.control <- trainControl(method="cv", number=5, allowParallel = TRUE)
model.gbm = train(y ~., method = "xgbTree", tuneGrid = grid.xg, data = khan.train, verbose = FALSE);
stopCluster(cluster)
registerDoSEQ()

plot(model.gbm)
```





```
predict.xgb <- predict.train(object=model.gbm, khan.test, type="raw")
print(head(predict.xgb))
```

```
## [1] 4 2 4 2 1 3
## Levels: 1 2 3 4
```

```
confusionMatrix(predict.xgb, khan.test$y)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction 1 2 3 4
```

```
##           1 3 0 0 0
```

```
##           2 0 6 0 0
```

```
##           3 0 0 5 0
```

```
##           4 0 0 1 5
```

```
##
```

```
## Overall Statistics
```

```
##
```

```
##           Accuracy : 0.95
```

```
##           95% CI : (0.7513, 0.9987)
```

```
##           No Information Rate : 0.3
```

```
##           P-Value [Acc > NIR] : 1.662e-09
```

```
##
```

```
##           Kappa : 0.9322
```

```
##           McNemar's Test P-Value : NA
```

```
##
```

```
## Statistics by Class:
##
##           Class: 1 Class: 2 Class: 3 Class: 4
## Sensitivity      1.00      1.0   0.8333   1.0000
## Specificity      1.00      1.0   1.0000   0.9333
## Pos Pred Value   1.00      1.0   1.0000   0.8333
## Neg Pred Value   1.00      1.0   0.9333   1.0000
## Prevalence       0.15      0.3   0.3000   0.2500
## Detection Rate   0.15      0.3   0.2500   0.2500
## Detection Prevalence 0.15      0.3   0.2500   0.3000
## Balanced Accuracy 1.00      1.0   0.9167   0.9667
```

```
#varImp(model.gbm)
```

Both models perform well, having test error rates of 5%. That said, the CI is smaller for the random forest model.

The random forest model a very wide search grid for mtry and resulted in a 5% error rate. This model was far harder to tune because with CV it takes much longer than XGB.

Chapter 9.8

(a)

```
set.seed(11)
train <- sample(1:nrow(OJ), 800)
oj.train <- OJ[train,]
oj.test <- OJ[-train,]
```

(b)

```
library(e1071)
svm.lin <- svm(Purchase ~ ., data = oj.train, kernel = "linear", cost = 0.01)
summary(svm.lin)
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = oj.train, kernel = "linear",
##      cost = 0.01)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##      cost:  0.01
##   gamma:  0.05555556
##
## Number of Support Vectors:  429
##
##   ( 216 213 )
##
##
## Number of Classes:  2
##
## Levels:
```

```
## CH MM
```

There were 429 support vectors, 216 for the class CH and 213 for the class MM.

(c)

```
svm.lin.pred <- predict(svm.lin, oj.train)
confusionMatrix(svm.lin.pred, oj.train$Purchase)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  CH  MM
##           CH 454  77
##           MM  49 220
##
##           Accuracy : 0.8425
##           95% CI : (0.8154, 0.8671)
##           No Information Rate : 0.6288
##           P-Value [Acc > NIR] : < 2e-16
##
##           Kappa : 0.656
##           McNemar's Test P-Value : 0.01616
##
##           Sensitivity : 0.9026
##           Specificity : 0.7407
##           Pos Pred Value : 0.8550
##           Neg Pred Value : 0.8178
##           Prevalence : 0.6288
##           Detection Rate : 0.5675
##           Detection Prevalence : 0.6637
##           Balanced Accuracy : 0.8217
##
##           'Positive' Class : CH
##
```

```
1-.8425
```

```
## [1] 0.1575
```

```
svm.lin.pred1 <- predict(svm.lin, oj.test)
confusionMatrix(svm.lin.pred1, oj.test$Purchase)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  CH  MM
##           CH 131  32
##           MM  19  88
##
##           Accuracy : 0.8111
##           95% CI : (0.7592, 0.856)
##           No Information Rate : 0.5556
##           P-Value [Acc > NIR] : < 2e-16
##
##           Kappa : 0.6133
##           McNemar's Test P-Value : 0.09289
```

```
##
##          Sensitivity : 0.8733
##          Specificity : 0.7333
##          Pos Pred Value : 0.8037
##          Neg Pred Value : 0.8224
##          Prevalence : 0.5556
##          Detection Rate : 0.4852
##          Detection Prevalence : 0.6037
##          Balanced Accuracy : 0.8033
##
##          'Positive' Class : CH
##
```

```
1-.8111
```

```
## [1] 0.1889
```

The training error rate is 15.75% and the test error rate is 18.89%.

(d)

```
set.seed(1)
tune.svm <- tune(svm, Purchase~., data = oj.train, kernel = "linear", ranges = list(cost = seq(.01, 10,
summary(tune.svm)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   0.41
##
## - best performance: 0.15875
##
## - Detailed performance results:
##   cost   error dispersion
## 1  0.01 0.16250 0.01666667
## 2  0.21 0.16125 0.02389938
## 3  0.41 0.15875 0.02045490
## 4  0.61 0.15875 0.02208726
## 5  0.81 0.16125 0.02389938
## 6  1.01 0.16375 0.02729087
## 7  1.21 0.16375 0.02729087
## 8  1.41 0.16500 0.02622022
## 9  1.61 0.16500 0.02622022
## 10 1.81 0.16500 0.02622022
## 11 2.01 0.16500 0.02622022
## 12 2.21 0.16500 0.02622022
## 13 2.41 0.16500 0.02622022
## 14 2.61 0.16500 0.02622022
## 15 2.81 0.16500 0.02622022
## 16 3.01 0.16375 0.02531057
## 17 3.21 0.16375 0.02531057
## 18 3.41 0.16375 0.02531057
## 19 3.61 0.16500 0.02415229
```

```
## 20 3.81 0.16500 0.02415229
## 21 4.01 0.16500 0.02415229
## 22 4.21 0.16500 0.02415229
## 23 4.41 0.16625 0.02503470
## 24 4.61 0.16625 0.02503470
## 25 4.81 0.16625 0.02503470
## 26 5.01 0.16625 0.02503470
## 27 5.21 0.16625 0.02503470
## 28 5.41 0.16625 0.02503470
## 29 5.61 0.16625 0.02503470
## 30 5.81 0.16625 0.02503470
## 31 6.01 0.16625 0.02503470
## 32 6.21 0.16625 0.02503470
## 33 6.41 0.16625 0.02503470
## 34 6.61 0.16625 0.02503470
## 35 6.81 0.16625 0.02503470
## 36 7.01 0.16625 0.02503470
## 37 7.21 0.16625 0.02503470
## 38 7.41 0.16625 0.02503470
## 39 7.61 0.16625 0.02503470
## 40 7.81 0.16625 0.02503470
## 41 8.01 0.16625 0.02503470
## 42 8.21 0.16625 0.02503470
## 43 8.41 0.16625 0.02503470
## 44 8.61 0.16625 0.02503470
## 45 8.81 0.16750 0.02443813
## 46 9.01 0.16750 0.02443813
## 47 9.21 0.16750 0.02443813
## 48 9.41 0.16750 0.02443813
## 49 9.61 0.16750 0.02443813
## 50 9.81 0.16625 0.02503470
```

The optimal cost = 0.41.

(e)

```
svm.lin <- svm(Purchase ~ ., data = oj.train, kernel = "linear", cost = .41)
summary(svm.lin)
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = oj.train, kernel = "linear",
##      cost = 0.41)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##      cost:   0.41
##   gamma:    0.05555556
##
## Number of Support Vectors: 326
##
## ( 161 165 )
##
##
```

```
## Number of Classes: 2
##
## Levels:
## CH MM
```

```
svm.lin.pred <- predict(svm.lin, oj.train)
confusionMatrix(svm.lin.pred, oj.train$Purchase)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction CH  MM
##           CH 453 76
##           MM  50 221
##
##           Accuracy : 0.8425
##           95% CI : (0.8154, 0.8671)
##           No Information Rate : 0.6288
##           P-Value [Acc > NIR] : < 2e-16
##
##           Kappa : 0.6565
##           McNemar's Test P-Value : 0.02594
##
##           Sensitivity : 0.9006
##           Specificity : 0.7441
##           Pos Pred Value : 0.8563
##           Neg Pred Value : 0.8155
##           Prevalence : 0.6288
##           Detection Rate : 0.5663
##           Detection Prevalence : 0.6613
##           Balanced Accuracy : 0.8224
##
##           'Positive' Class : CH
##
```

```
1-.8425
```

```
## [1] 0.1575
```

```
svm.lin.pred1 <- predict(svm.lin, oj.test)
confusionMatrix(svm.lin.pred1, oj.test$Purchase)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction CH  MM
##           CH 130 33
##           MM  20 87
##
##           Accuracy : 0.8037
##           95% CI : (0.7512, 0.8494)
##           No Information Rate : 0.5556
##           P-Value [Acc > NIR] : < 2e-16
##
##           Kappa : 0.5981
##           McNemar's Test P-Value : 0.09929
```



```
##
##          Sensitivity : 0.8667
##          Specificity : 0.7250
##          Pos Pred Value : 0.7975
##          Neg Pred Value : 0.8131
##          Prevalence : 0.5556
##          Detection Rate : 0.4815
##          Detection Prevalence : 0.6037
##          Balanced Accuracy : 0.7958
##
##          'Positive' Class : CH
##
```

```
1-.8037
```

```
## [1] 0.1963
```

Training error = 15.75%, test error = 19.63%. Not an improvement.

(f)

```
svm.rad <- svm(Purchase ~ ., data = oj.train, kernel = "radial")
summary(svm.rad)
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = oj.train, kernel = "radial")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##         cost:  1
##        gamma: 0.05555556
##
## Number of Support Vectors: 362
##
## ( 182 180 )
##
##
## Number of Classes: 2
##
## Levels:
##  CH MM
```

```
svm.rad.pred <- predict(svm.rad, oj.train)
confusionMatrix(svm.rad.pred, oj.train$Purchase)
```

```
## Confusion Matrix and Statistics
##
##          Reference
## Prediction  CH  MM
##      CH 463   80
##      MM  40 217
##
##          Accuracy : 0.85
##          95% CI : (0.8233, 0.874)
```

```
##      No Information Rate : 0.6288
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.6696
##  Mcnemar's Test P-Value : 0.0003706
##
##      Sensitivity : 0.9205
##      Specificity : 0.7306
##      Pos Pred Value : 0.8527
##      Neg Pred Value : 0.8444
##      Prevalence : 0.6288
##      Detection Rate : 0.5787
##      Detection Prevalence : 0.6787
##      Balanced Accuracy : 0.8256
##
##      'Positive' Class : CH
##
```

1-.85

```
## [1] 0.15
```

```
svm.rad.pred1 <- predict(svm.rad, oj.test)
confusionMatrix(svm.rad.pred1, oj.test$Purchase)
```

```
## Confusion Matrix and Statistics
##
##      Reference
## Prediction  CH  MM
##      CH 133  34
##      MM  17  86
##
##      Accuracy : 0.8111
##      95% CI : (0.7592, 0.856)
##      No Information Rate : 0.5556
##      P-Value [Acc > NIR] : < 2e-16
##
##              Kappa : 0.612
##  Mcnemar's Test P-Value : 0.02506
##
##      Sensitivity : 0.8867
##      Specificity : 0.7167
##      Pos Pred Value : 0.7964
##      Neg Pred Value : 0.8350
##      Prevalence : 0.5556
##      Detection Rate : 0.4926
##      Detection Prevalence : 0.6185
##      Balanced Accuracy : 0.8017
##
##      'Positive' Class : CH
##
```

1-.8111

```
## [1] 0.1889
```

Radial SVM train error is 15%, test 18.89%. This is not an improvement over linear SVM.

```
set.seed(1)
tune.svm <- tune(svm, Purchase~., data = oj.train, kernel = "radial", ranges = list(cost = seq(.01, 10,
summary(tune.svm)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   0.61
##
## - best performance: 0.16625
##
## - Detailed performance results:
##   cost   error dispersion
## 1  0.01 0.37125 0.04752558
## 2  0.21 0.17500 0.01954340
## 3  0.41 0.17000 0.02776389
## 4  0.61 0.16625 0.02503470
## 5  0.81 0.16625 0.02703521
## 6  1.01 0.16625 0.02949223
## 7  1.21 0.16875 0.02447363
## 8  1.41 0.16875 0.02585349
## 9  1.61 0.17000 0.03184162
## 10 1.81 0.16875 0.02960973
## 11 2.01 0.17000 0.02958040
## 12 2.21 0.17250 0.02813657
## 13 2.41 0.17500 0.02886751
## 14 2.61 0.17500 0.02886751
## 15 2.81 0.17500 0.02886751
## 16 3.01 0.17875 0.03175973
## 17 3.21 0.18000 0.03238227
## 18 3.41 0.18000 0.03073181
## 19 3.61 0.17875 0.03064696
## 20 3.81 0.17875 0.03064696
## 21 4.01 0.17875 0.03064696
## 22 4.21 0.17875 0.03064696
## 23 4.41 0.17875 0.03064696
## 24 4.61 0.17875 0.03064696
## 25 4.81 0.17875 0.02829041
## 26 5.01 0.17750 0.02874698
## 27 5.21 0.17750 0.02874698
## 28 5.41 0.17750 0.02874698
## 29 5.61 0.17750 0.02874698
## 30 5.81 0.17875 0.03007514
## 31 6.01 0.17750 0.02751262
## 32 6.21 0.17750 0.02751262
## 33 6.41 0.17875 0.02889757
## 34 6.61 0.17750 0.03050501
## 35 6.81 0.18000 0.02838231
## 36 7.01 0.18000 0.02838231
## 37 7.21 0.18125 0.03076005
```

```
## 38 7.41 0.18250 0.03184162
## 39 7.61 0.18250 0.03184162
## 40 7.81 0.18375 0.03064696
## 41 8.01 0.18250 0.02898755
## 42 8.21 0.18250 0.02898755
## 43 8.41 0.18250 0.02838231
## 44 8.61 0.18250 0.02838231
## 45 8.81 0.18250 0.02838231
## 46 9.01 0.18250 0.02838231
## 47 9.21 0.18375 0.02829041
## 48 9.41 0.18375 0.02829041
## 49 9.61 0.18375 0.02829041
## 50 9.81 0.18375 0.02829041
```

```
svm.rad <- svm(Purchase ~ ., data = oj.train, kernel = "radial", cost = .61)
summary(svm.rad)
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = oj.train, kernel = "radial",
##      cost = 0.61)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##         cost: 0.61
##        gamma: 0.05555556
##
## Number of Support Vectors: 384
##
## ( 193 191 )
##
##
## Number of Classes: 2
##
## Levels:
## CH MM
```

```
svm.rad.pred <- predict(svm.rad, oj.train)
confusionMatrix(svm.rad.pred, oj.train$Purchase)
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction  CH  MM
##          CH 462  77
##          MM  41 220
##
##              Accuracy : 0.8525
##              95% CI : (0.826, 0.8764)
##      No Information Rate : 0.6288
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.676
```

```
## McNemar's Test P-Value : 0.001273
##
##      Sensitivity : 0.9185
##      Specificity : 0.7407
##      Pos Pred Value : 0.8571
##      Neg Pred Value : 0.8429
##      Prevalence : 0.6288
##      Detection Rate : 0.5775
##      Detection Prevalence : 0.6737
##      Balanced Accuracy : 0.8296
##
##      'Positive' Class : CH
##
```

```
1-.8525
```

```
## [1] 0.1475
```

```
svm.rad.pred1 <- predict(svm.rad, oj.test)
confusionMatrix(svm.rad.pred1, oj.test$Purchase)
```

```
## Confusion Matrix and Statistics
##
##      Reference
## Prediction  CH  MM
##      CH 133  36
##      MM  17  84
##
##      Accuracy : 0.8037
##      95% CI : (0.7512, 0.8494)
##      No Information Rate : 0.5556
##      P-Value [Acc > NIR] : < 2e-16
##
##      Kappa : 0.5961
##      McNemar's Test P-Value : 0.01342
##
##      Sensitivity : 0.8867
##      Specificity : 0.7000
##      Pos Pred Value : 0.7870
##      Neg Pred Value : 0.8317
##      Prevalence : 0.5556
##      Detection Rate : 0.4926
##      Detection Prevalence : 0.6259
##      Balanced Accuracy : 0.7933
##
##      'Positive' Class : CH
##
```

```
1-.8037
```

```
## [1] 0.1963
```

Again, not an improvement: tune—or at least the way I am using it—seems to increase variance.

(g)

```
svm.poly<- svm(Purchase ~ ., data = oj.train, kernel = "polynomial", degree = 2)
summary(svm.poly)
```

```

##
## Call:
## svm(formula = Purchase ~ ., data = oj.train, kernel = "polynomial",
##      degree = 2)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: polynomial
##      cost:   1
##   degree:   2
##   gamma:   0.05555556
##   coef.0:   0
##
## Number of Support Vectors:  425
##
## ( 215 210 )
##
##
## Number of Classes:  2
##
## Levels:
##  CH MM

svm.poly.pred <- predict(svm.poly, oj.train)
confusionMatrix(svm.poly.pred, oj.train$Purchase)

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  CH  MM
##           CH 471 103
##           MM  32 194
##
##           Accuracy : 0.8312
##           95% CI : (0.8035, 0.8566)
##   No Information Rate : 0.6288
##   P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.6199
##   McNemar's Test P-Value : 1.695e-09
##
##           Sensitivity : 0.9364
##           Specificity : 0.6532
##           Pos Pred Value : 0.8206
##           Neg Pred Value : 0.8584
##           Prevalence : 0.6288
##           Detection Rate : 0.5887
##   Detection Prevalence : 0.7175
##           Balanced Accuracy : 0.7948
##
##           'Positive' Class : CH
##

```

```
1-.8312
```

```
## [1] 0.1688
```

```
svm.poly.pred1 <- predict(svm.poly, oj.test)
confusionMatrix(svm.poly.pred1, oj.test$Purchase)
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction  CH  MM
##          CH 137  49
##          MM  13  71
##
##              Accuracy : 0.7704
##              95% CI : (0.7155, 0.8192)
##      No Information Rate : 0.5556
##      P-Value [Acc > NIR] : 1.621e-13
##
##              Kappa : 0.5206
##  Mcnemar's Test P-Value : 8.789e-06
##
##              Sensitivity : 0.9133
##              Specificity : 0.5917
##              Pos Pred Value : 0.7366
##              Neg Pred Value : 0.8452
##              Prevalence : 0.5556
##              Detection Rate : 0.5074
##      Detection Prevalence : 0.6889
##              Balanced Accuracy : 0.7525
##
##              'Positive' Class : CH
##
```

```
1-.7704
```

```
## [1] 0.2296
```

The polynomial SVM of degree 2 has a train error of 16.88% and a test error of 22.96%, markedly worse than the linear SVM.

```
set.seed(1)
tune.svm <- tune(svm, Purchase~., data = oj.train, kernel = "polynomial", degree = 2, ranges = list(cost = 100))
summary(tune.svm)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   7.01
##
## - best performance: 0.165
##
## - Detailed performance results:
```

```
##      cost   error dispersion
## 1  0.01 0.37250 0.04779877
## 2  0.21 0.23125 0.04611655
## 3  0.41 0.19375 0.03547789
## 4  0.61 0.19500 0.03129164
## 5  0.81 0.18750 0.03004626
## 6  1.01 0.19000 0.03322900
## 7  1.21 0.18500 0.03763863
## 8  1.41 0.18125 0.03346329
## 9  1.61 0.18125 0.02960973
## 10 1.81 0.18125 0.02716334
## 11 2.01 0.18250 0.02776389
## 12 2.21 0.18375 0.02949223
## 13 2.41 0.18250 0.03291403
## 14 2.61 0.18375 0.03387579
## 15 2.81 0.18250 0.03496029
## 16 3.01 0.18250 0.03446012
## 17 3.21 0.18000 0.03736085
## 18 3.41 0.18125 0.03596391
## 19 3.61 0.17750 0.03670453
## 20 3.81 0.17750 0.03670453
## 21 4.01 0.17625 0.03557562
## 22 4.21 0.17500 0.03632416
## 23 4.41 0.17625 0.03928617
## 24 4.61 0.17500 0.03818813
## 25 4.81 0.17500 0.04082483
## 26 5.01 0.17375 0.03884174
## 27 5.21 0.17250 0.03717451
## 28 5.41 0.17250 0.03717451
## 29 5.61 0.16875 0.03547789
## 30 5.81 0.17000 0.03446012
## 31 6.01 0.17000 0.03446012
## 32 6.21 0.17125 0.03488573
## 33 6.41 0.16750 0.03782269
## 34 6.61 0.16875 0.03784563
## 35 6.81 0.16625 0.03682259
## 36 7.01 0.16500 0.03622844
## 37 7.21 0.16500 0.03574602
## 38 7.41 0.16500 0.03574602
## 39 7.61 0.16750 0.03129164
## 40 7.81 0.16750 0.03129164
## 41 8.01 0.16875 0.03240906
## 42 8.21 0.16875 0.03240906
## 43 8.41 0.16875 0.03240906
## 44 8.61 0.16875 0.03240906
## 45 8.81 0.16875 0.03240906
## 46 9.01 0.16750 0.03238227
## 47 9.21 0.16750 0.03016160
## 48 9.41 0.16625 0.03230175
## 49 9.61 0.16625 0.03230175
## 50 9.81 0.16625 0.03230175
```

```
svm.poly<- svm(Purchase ~ ., data = oj.train, kernel = "polynomial", degree = 2, cost = 7.01)
summary(svm.poly)
```



```

##
## Call:
## svm(formula = Purchase ~ ., data = oj.train, kernel = "polynomial",
##      degree = 2, cost = 7.01)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: polynomial
##      cost:   7.01
##   degree:    2
##   gamma:    0.05555556
##   coef.0:    0
##
## Number of Support Vectors:  337
##
##   ( 170 167 )
##
##
## Number of Classes:  2
##
## Levels:
##   CH MM

```

```

svm.poly.pred <- predict(svm.poly, oj.train)
confusionMatrix(svm.poly.pred, oj.train$Purchase)

```

```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  CH  MM
##           CH 467  82
##           MM  36 215
##
##           Accuracy : 0.8525
##           95% CI : (0.826, 0.8764)
##   No Information Rate : 0.6288
##   P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.6737
##   McNemar's Test P-Value : 3.434e-05
##
##           Sensitivity : 0.9284
##           Specificity : 0.7239
##   Pos Pred Value : 0.8506
##   Neg Pred Value : 0.8566
##           Prevalence : 0.6288
##   Detection Rate : 0.5837
##   Detection Prevalence : 0.6863
##   Balanced Accuracy : 0.8262
##
##   'Positive' Class : CH
##

```

```
1-.8525
```

```
## [1] 0.1475
```

```
svm.poly.pred1 <- predict(svm.poly, oj.test)
confusionMatrix(svm.poly.pred1, oj.test$Purchase)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  CH  MM
##           CH 136  40
##           MM  14  80
##
##           Accuracy : 0.8
##           95% CI : (0.7472, 0.846)
##           No Information Rate : 0.5556
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.586
##           McNemar's Test P-Value : 0.0006688
##
##           Sensitivity : 0.9067
##           Specificity : 0.6667
##           Pos Pred Value : 0.7727
##           Neg Pred Value : 0.8511
##           Prevalence : 0.5556
##           Detection Rate : 0.5037
##           Detection Prevalence : 0.6519
##           Balanced Accuracy : 0.7867
##
##           'Positive' Class : CH
##
```

```
1-.8
```

```
## [1] 0.2
```

Using the tuned cost parameter, the train error decreases to 14.75% and the test error to 20%.

(h)

The linear SVM when cost = .01 or the radial SVM with the default settings give the same performance.